



Formalization and Runtime Verification of Invariants for Robotic Systems

Ricardo Jorge Dias Cordeiro

Mestrado em Engenharia Informática
Especialização em Interação e Conhecimento

Dissertação orientada por:
Prof. Doutor Alcides Miguel Cachulo Aguiar Fonseca
Prof. Doutor Christopher Steven Timperley

Acknowledgments

I would like to thank my coordinator, Prof. Alcides Fonseca, for the exceptional way of teaching not only through the making of my thesis but also throughout all my academic courses.

My coordinator Prof. Chris Timperley for taking his time to help me in this chapter of his life where he had to take care of his baby.

My upperclassman Paulo and Catarina, for all the advice and help.

All my friends who spent time with me know that somehow you helped me through this process.

All my family, in particular my grandparents.

My little brother.

In the end, and more importantly, my mother for taking care of me all my life and giving me the opportunity to follow this path.

"Dreams breathe life into men, and can cage them in suffering. Men live and die by their dreams, but long after they've been abandoned, they still smolder deep in men's hearts. Some see nothing more than life and death. They are dead! For they have no dreams."

Kentaro Miura in Berserk

Resumo

A Robótica tem uma grande influência na sociedade atual, seja industrialmente, na medicina, na agricultura, ou como forma de lazer, e, muitas vezes, toma um papel crucial em que uma falha em algum destes sistemas robóticos cruciais pode impactar o modo em como nós vivemos. Se, por exemplo, um carro autónomo provocar a morte de algum passageiro devido a um defeito, futuros e atuais utilizadores deste modelo irão certamente ficar apreensivos em relação à sua utilização. Assegurar que robôs reproduzam um comportamento correto pode assim salvar bastante dinheiro em estragos ou até mesmo as nossas vidas.

Os sistemas robóticos são não determinísticos, isto porque os robôs interagem diretamente com o mundo real. Testar *software* num destes ambientes é bastante complexo devido às variáveis serem imprevisíveis e mudarem constantemente. Verificar o sucesso de um movimento ou tarefa neste ambiente pode ser bastante difícil do ponto de vista de um robô, pelo que monitorização externa é muitas vezes necessária.

Devido à ampla utilização de sistemas robóticos, a qualidade do software que corre em robôs deveria ser de extrema importância para nós. O *software* destes sistemas, assim como os métodos utilizados para os testar, são bastante específicos da área e diferentes de *software* tradicional, em grande parte devido à sua já falada interação com o mundo real.

As práticas atuais em relação à testagem de sistemas robóticos são vastas e envolvem métodos como simulações, verificação de “logs”, testes em campo, entre outras. Frequentemente, um denominador comum entre as práticas adotadas é a necessidade de um humano pessoalmente analisar e determinar se o comportamento de um sistema robótico é o correto. A automatização deste tipo de análise poderia não só aliviar o trabalho de técnicos especializados, facilitando assim a realização de testes, mas também possivelmente permitir a execução massiva de testes em paralelo. Este tipo de testagem automática poderá potencialmente detetar falhas no comportamento do sistema robótico que de outra maneira não seriam identificadas devido a erros humanos ou mesmo à falta de tempo.

Uma invariante representa uma propriedade que se mantém durante toda a execução do sistema, dispor de uma lista de invariantes para um sistema robótico e ser capaz de as verificar em tempo de execução é uma forma de provar a qualidade desse sistema.

Atualmente já existe investigação e literatura substancial relacionada com a utilização de testes automáticos em sistemas robóticos, assim como ferramentas para realizar de alguma maneira este tipo de análise. No entanto, de uma forma geral, a automatização no campo da deteção de erros ou até mesmo na utilização de invariantes continua a não ser adotada para este tipo de sistemas. O problema na adoção deste tipo de ferramentas deve-se à complexidade ou à

falta de confiança nas soluções já desenvolvidas, algo que incentiva o estudo apresentado nesta tese.

Esta dissertação visa assim explorar o problema da automatização da detecção de erros comportamentais em robôs num ambiente de simulação, introduzindo uma linguagem de domínio específico direcionada a especificar as propriedades de sistemas robóticos em relação ao seu ambiente, assim como a geração de *software* de monitorização capaz de detetar a transgressão destas propriedades durante uma simulação.

A linguagem de domínio específico também assume que o sistema robótico irá ser executado por meio da framework de código aberto ROS (Robot Operating System). O ROS é uma framework que oferece uma vasta coleção de bibliotecas, interfaces e ferramentas especificamente desenhadas para ajudar na construção de *software* para robôs. O ROS fornece uma abstração entre *software* e hardware que ajuda desenvolvedores a facilmente conectar diferentes componentes de robôs através de mensagens enviadas por canais de comunicação chamados *tópicos*. O ROS tem uma arquitetura modular e outras vantagens que têm como objectivo a intra-colaboração e fácil desenvolvimento de *software*. O seu ecossistema está construído de maneira a que a maioria dos projetos dependem de uma pequena lista de pacotes. Devido a todos os factos em cima mencionadas o ROS é amplamente utilizado para investigação e na indústria da robótica, e por essas mesmas razões o escolhemos como a framework que seria integrada neste trabalho.

A Lógica Temporal Linear pode ser usada como um método de verificação de programas e também ser útil para a criação de invariantes em sistemas robóticos. Um sistema formal de lógica temporal contém padrões que podem ser usados como uma forma de especificação de propriedades deste tipo de sistemas.

A Lógica Temporal Linear é um ramo da lógica responsável por representar e relacionar componentes em referência a uma linha temporal. A linguagem de domínio específico necessita de expressar requisitos de determinados estados ou eventos durante a simulação, desta maneira precisa de apresentar determinadas características: Palavras-chave para representar relações temporais de ou entre objetos, como, por exemplo, o robô "nunca", ou "eventualmente" o robô. Referências a estados anteriores da simulação, como, por exemplo, a velocidade do robô está sempre a aumentar, ou seja, é sempre maior que no estado anterior. Atalhos para ser possível referir certas características de ou entre objetos, como, por exemplo, a "posição", "velocidade" ou "distância" de ou entre robôs.

O *software* de monitorização gerado é um ficheiro *Python* que tem origem numa especificação feita através da linguagem de domínio específico e que correrá sobre a framework ROS. A geração deste ficheiro assume também que a monitorização será feita no simulador Gazebo, isto porque para obter dados como a posição ou velocidade absoluta de um robô durante a simulação é necessário aceder a *tópicos* ROS específicos que na geração do ficheiro de monitorização estão *hardcoded*, pois dependem do *software* de monitorização escolhido.

O Gazebo é um simulador de alta-fidelidade capaz de simular sistemas robóticos em qualquer tipo de ambiente ou condições. O Gazebo é um simulador de código aberto que suporta ferramentas como a simulação de sensores, manipulação de modelos, e o controlo de atuadores sob diferentes motores de física, entre outras ferramentas, o que o faz um simulador que vários

sistemas robóticos diferentes entre si são capazes de tirar proveito, daí a sua escolha para o desenvolvimento deste trabalho.

A geração de um ficheiro capaz de executar a monitorização durante uma simulação significa também que este tipo de monitorização pode ser executada independente de um sistema robótico, permitindo assim a automatização da monitorização a respeito de vários objetos e as suas relações.

O objetivo deste trabalho é então fornecer uma maneira de desenvolvedores de *software* para sistemas robóticos de conseguirem verificar propriedades posicionais e temporais dos seus sistemas de maneira automática através de uma linguagem de domínio específico, que deve, ao mesmo tempo, ser simples, intuitiva e permitir expressar propriedades que sejam relevantes entre componentes da simulação.

De maneira a avaliar o trabalho desenvolvido tentamos detetar erros em sistemas robóticos, inserindo propositadamente *bugs* no sistema. Estes *bugs* provêm de uma lista de *bugs* que foram previamente identificados por outros utilizadores destes sistemas robóticos. A lista contém *bugs* bastante diferentes entre si, dos quais a maior parte já foi corrigido em *software* atual, sendo que o nosso objetivo é identificar *bugs* antigos que ocorram em tempo de execução.

Resultados mostram que é possível expressar propriedades temporais e posicionais de e entre robôs e o seu ambiente com o suporte da linguagem de domínio específico. O trabalho mostra também que é possível automatizar a monitorização da violação de alguns tipos de comportamentos esperados de robôs em relação ao seu estado ou determinados eventos que ocorrem durante uma simulação.

Palavras-chave: Robótica, Linguagem de domínio, Monitorização em tempo de execução, Detecção de erros

Abstract

Robotic systems are critical in today's society, be it in manufacturing, medicine, or agriculture. A potential failure in a robot may have extraordinary costs, not only financial but can also cost lives.

Current practices in robot testing are vast and involve methods like simulation, log checking, or field testing. However, current practices often require human monitoring to determine the correctness of a given behavior. Automating this analysis can not only relieve the burden from a high-skilled engineer but also allow for massive parallel executions of tests that can detect behavioral faults in the robots. These faults could otherwise not be found due to human error or a lack of time.

I have developed a Domain Specific Language to specify the properties of robotic systems in the Robot Operating System (ROS). Developer written specifications in this language compile to a monitor ROS module that detects violations of those properties in runtime. I have used this language to express the temporal and positional properties of robots using Linear Temporal Logic as a basis for the language stipulation. I have also automated the monitoring of some behavioral violations of robots in relation to their state or events during a simulation, resorting to relations between the internal information of the system and the corresponding information in the simulator.

To evaluate the developed work, I went through a list of documented ROS bugs and identified some that happen at runtime. Using these bugs as a basis I specified the robot's properties in the developed language that should be capable of detecting an error, in order to test both the expressiveness and the monitoring while running the system.

Keywords: Robotics, Domain-specific language, Runtime Monitoring, Error detection

Contents

| | |
|---|-----------|
| List of Figures | xv |
| 1 Introduction | 1 |
| 1.1 Motivation | 1 |
| 1.2 Problem Statement | 2 |
| 1.3 Objectives | 2 |
| 1.4 Motivational Example | 3 |
| 1.5 Contributions | 3 |
| 1.6 Structure of the document | 4 |
| 2 Background & Related Work | 5 |
| 2.1 Background | 5 |
| 2.1.1 Robot Operating System | 5 |
| 2.1.2 Gazebo | 5 |
| 2.2 Related Work | 6 |
| 2.2.1 Runtime Testing | 6 |
| 2.2.2 Linear Temporal Logic and Invariants | 6 |
| 2.2.3 Monitoring Frameworks | 7 |
| 3 A Specification Language for Robotics Properties | 9 |
| 3.1 Language Notations | 9 |
| 3.1.1 High-level Concepts | 9 |
| 3.1.2 Temporal Keywords | 9 |
| 3.1.3 Temporal value | 10 |
| 3.1.4 Simulation primitives | 10 |
| 3.1.5 Operands | 10 |
| 3.1.6 Operators | 11 |
| 3.1.7 Protected Variables | 11 |
| 3.1.8 Topic declaration | 11 |
| 3.1.9 Model robots | 11 |
| 3.2 Grammar | 11 |
| 3.3 DSL Usage Examples | 13 |
| 3.3.1 Vehicle Maximum Speed | 13 |

| | | |
|----------|--|-----------|
| 3.3.2 | Follow the Leader | 13 |
| 3.3.3 | Localization error | 13 |
| 3.3.4 | Drone height rotors control | 13 |
| 4 | Monitoring | 15 |
| 4.1 | Runtime Monitoring | 15 |
| 4.2 | Generated File | 15 |
| 4.3 | Error Messages | 16 |
| 5 | Evaluation | 19 |
| 5.1 | Evaluation Overview | 19 |
| 5.2 | Experiments | 20 |
| 5.2.1 | Calculation Error Inverts Turning Direction | 20 |
| 5.2.2 | Robot Getting Stuck When Auto-docking | 21 |
| 5.2.3 | Unexpected Movement Due to Wrong Calculation | 22 |
| 6 | Future Work | 25 |
| 6.1 | Performance Tweaking | 25 |
| 6.2 | Validation of the Proposal | 25 |
| 6.3 | Better Error Messages | 25 |
| 6.4 | Integrate the DSL with Scenario Generation Tools | 26 |
| 6.5 | Integration with other Industrial Simulators | 26 |
| 7 | Conclusion | 27 |
| | References | 30 |

List of Figures

| | | |
|-----|---|----|
| 1.1 | Example of the displayed error when the robot does not stop at the stop sign. . . | 3 |
| 4.1 | Example of an error message. | 17 |
| 5.1 | The normal runtime flow of the system. | 20 |
| 5.2 | The flow of the system when the adding the daemon node. | 20 |
| 5.3 | Calculation Error Inverts Turning Direction bug error message. | 21 |
| 5.4 | Robot Getting Stuck When Auto-docking bug error message. | 22 |
| 5.5 | Unexpected Movement Due to Wrong Calculation bug error message. | 23 |

Chapter 1

Introduction

This thesis explores a possible solution for automation in the testing of robotic systems through a domain-specific language (DSL) and simulation-based monitoring software.

A DSL is a programming language that is written for a specific domain, it usually provides a higher level of abstraction and is simpler than other languages mainly because they are intended to be used by people knowledgeable of the domain.

This chapter intends to introduce the motivation for this work (section 1.1), present the problem statements of such an approach (section 1.2), discuss the objectives (section 1.3), show a motivational example of the developed work (section 1.4), present the expected contributions (section 1.5), and finally summarize the structure of the rest of the document (section 1.6).

1.1 Motivation

Robotics already significantly impact our current society, in industry, in medicine, in agriculture, or leisurely in sports contests or personal use. Robotics often take critical roles like the example of robot arms in car assembly lines or autonomous farms. The tendency is for robot usage to keep growing at a global level.

Robotic Systems are non-deterministic, mainly because robots interact directly with the real world. Testing software in such environments is complex, as many variables can change, and verifying the success of a task or movement may not be possible from the robot's perspective, and external monitoring may be required.

Current practices in testing robot software mainly involve field testing, simulation testing, and log checking and require a human to analyze the robot's behavior to determine whether the behavior is correct. Due to their broad practicality, the quality of software running on robots should be extremely important to us. Robot software, as well as the techniques used to test their quality, are very field-specific and different from the techniques employed in traditional Software Engineering mainly because of their real-world interaction. This peculiarity means automatic tests are rarely used in robotics [10, 14].

Studying possible options for viable automation of tests in robotic systems could lead to an opening on its usage in both research and the industry. Also, allowing for multiple parallel executions of tests not depending on human monitoring could improve the quality of current and future robot software.

1.2 Problem Statement

The multiple challenges in robot testing influence the planning for testing a robot because there are tradeoffs among choices.

Using simulation-based testing, developers can take advantage of real values of objects' attributes to compare with what the robot system perceives. Using this alleyway, it is possible to, in a way, surpassing the need for human-in-the-loop testing.

Human-in-the-loop testing refers to a method of testing that requires a human to perform a certain task that can't be covered by means of automation or simulation.

While simulation-based tests are a promising approach for automation, there is still distrust in the precision and validity of the results. As a result, simulation-based autonomous testing is rarely used due to reliability and factors like cost and complexity [10, 14]. Due to these factors, despite being dangerous, sometimes expensive, or work-intensive, real-life robot testing or other methods are still the main choices. The resulting product is a lack of quality in the software across projects.

In this thesis, I address problems of defining simulation-based automated tests for robotics systems.

1.3 Objectives

The ultimate goal of this thesis is to remove the need for human-in-the-loop testing of robotic systems by studying a possible solution for automation in simulation-based tests.

This work aims to provide developers with a way to verify their robotic systems' properties in relation to their position in a simulation (positional properties) aswell as correlations between current and past events (temporal properties). To this end, I propose the introduction of a DSL for developers to express their relevant properties. The given properties compile into monitors that can be used in simulation to ensure the correctness of the system. The DSL was designed from the point of view of the Robot Operating System (ROS) [12] developers and tries to abstract the underlying Linear Temporal Logic (LTL) system. LTL is a branch of logic responsible for representing and reasoning about modalities in reference to time. The DSL allows properties to reason about native ROS constructs, like *topics*, *messages* and simulation information. Thus, it is possible to express properties that relate the internal information of the system with the corresponding information in the simulator.

The DSL should allow describing a robotic system's properties simply and intuitively while simultaneously expressing relevant temporal and positional arguments between robots components and objects in the simulation. For this reason, the expected design goals of the DSL that I believe conform with ROS developers are:

- Temporal operators based on but not restricted to Linear Temporal Logic.
- Primitives to make references to the simulation environment, like the position, velocity, or any others.

```

Error at line 3:
  after_until robot.distance.stop_sign < 1, robot.distance.stop_sign > 1, eventually
robot.velocity == 0
Failing state:
  robot.distance.stop_sign: 1.000545118597548
  robot.velocity: 0.17758309727799252

```

Figure 1.1: Example of the displayed error when the robot does not stop at the stop sign.

- Basic operators to make comparisons between defined components, like greater than, equals, and others.

1.4 Motivational Example

Let us consider an autonomous car developer wanting to express that the car always stops when near a stop sign. The following example presents a property defined in the language that specifies the intended behavior of the developer.

```
after_until robot.distance.stop_sign < 1, robot.distance.stop_sign > 1, eventually robot.velocity == 0
```

Translating into natural language, the property states in the first section that after the robot's distance to the stop-sign is below the value of 1 in the simulator, and in the second section that up until the distance is again above 1, then in the third section the robot velocity will eventually be equal to 0.

The toolchain compiles the DSL specification to executable Python code that is capable of running as a ROS node. The node listens only to relevant topics and performs the computations to verify the specified property.

The flow of the process of monitoring a robotic system is described as follows:

- (i) **Property formalization:** The developer describes in the DSL the properties of the robotic system one wants to monitor in a text file.
- (ii) **Compilation:** The specified properties are compiled, and a Python file, capable of running as a ROS node, is generated.
- (iii) **Monitoring:** The node can be run whenever testing the system and will listen to pertinent topics and perform the computations needed to verify the specified properties.

1.5 Contributions

The expected contributions of this thesis are below enumerated.

1. Definition of a domain-specific language to specify robotic systems' properties.
2. Implementation of a compiler for the language that can generate software capable of monitoring relevant components while in a simulation.
3. Evaluation of the expressive capabilities of the solution.

1.6 Structure of the document

The document is organized as follows:

- **Chapter 2** - Background & Related Work:
- **Chapter 3** - Specification Language for Robotics Properties
- **Chapter 4** - Monitoring
- **Chapter 5** - Evaluation
- **Chapter 6** - Future Work
- **Chapter 7** - Conclusion

Chapter 2

Background & Related Work

This chapter gives an overview of the background software adopted while developing this work (section 2.1), and shines light on the already existing similar work and adopted techniques on the subject (section 2.2).

2.1 Background

This section provides some background on the used software and the reason for its choice, what the Robot Operating System is (subsection 2.1.1), and the simulation software adopted (subsection 2.1.2).

2.1.1 Robot Operating System

The Robot Operating System (ROS) [12] is an open-source framework with a vast collection of libraries, interfaces, and tools designed to help build robot software. ROS provides an abstraction between hardware and software that helps developers easily connect the different robot components predominantly through messages sent through communication channels called *topics* via a publish-subscribe architecture.

ROS has a modular architecture and other advantages built with the purpose of cross-collaboration and easy development [3]. For all these reasons, ROS is used by hundreds of companies and research labs.

The ROS ecosystem ended up with a very strong «standard library» of packages that are used by almost everyone, and a large number of packages that don't belong to that standard library that are rarely used by anyone. This is generally seen as a bad thing in software ecosystems, but for this work, it is an advantage because I can specify the behavior of a smaller standard library and still be able to represent the behavior of a large number of ROS systems. Literature states that around eighty-two percent of ROS applications rely on packages released by a small subset of groups [10].

2.1.2 Gazebo

Robotic systems simulation is an essential tool for testing robots' behavior. For this reason, Gazebo [9] started with the idea of a high-fidelity simulator to simulate robots in any environment under mixed conditions.

Gazebo is an open-source 3D simulator that supports tools like sensors simulation, mesh management, and actuators control under different physics engines, among others, which makes it a simulator that very distinct robotic systems can use.

2.2 Related Work

This section shows some research on runtime testing, the different techniques, and the difficulties of implementing it that already exist (subsection 2.2.1). The importance of Invariant specification and its relation to Linear Temporal Logic (LTL) (subsection 2.2.2). In the end, some monitoring frameworks have already tried to implement similar runtime verification concepts (subsection 2.2.3).

2.2.1 Runtime Testing

Runtime Testing is an analysis that takes place during the execution of a program, it takes advantage of taking information from the running system to make inferences on if the observed information violates certain properties of the program.

Due to the mentioned unforeseen circumstances when executing robotic systems, runtime testing, although sometimes time-consuming, may be advantageous when identifying errors in these types of systems.

Implementing runtime monitoring adds load to the simulation. Therefore, not demanding excessive resources is essential when taking this approach.

Some challenges in implementing such mechanisms are mentioned in the cited paper [13].

Todo >Paper cited where? This paper needs proper introduction and relevance explained. What are the challenges?<

Besides the method of monitoring chosen in this work, runtime testing can be implemented in other ways. For instance, Mithra [5] is a tool that provides an oracle for automated simulation-based testing relying on machine learning software, more specifically, a three-step multivariate time series clustering. **Todo** >Seems out of nowhere. I don't see the relationship to your work from this paragraph. What other methods are there?<

2.2.2 Linear Temporal Logic and Invariants

An invariant represents a property that holds through the execution of the system. Having a set of invariants for a robotic system and asserting them at runtime makes it able to prove the correctness of the system.

Research on invariant checking [14] demonstrates that important safety bugs in real-world autonomous robotic systems can be identified when representing safety violations of systems and monitoring them.

Linear temporal logic (LTL) is a branch of logic responsible for representing and reasoning about modalities in reference to time.

As an approach for program verification, a formal system of temporal logic was suggested for both sequential and parallel programs [11]. LTL can be used as a method of model-checking [6]

using its patterns as a form of property specification. It includes patterns such as "always", "finally", "until", "eventually", and others, which can be useful in the creation of invariants for program verification of robotics systems.

2.2.3 Monitoring Frameworks

Similar work on runtime monitoring that integrates with ROS already exists.

ROSMonitoring [7] can monitor and log errors at the level of *topic* malfunctioning, however, it was designed with portability and scalability in mind which means it is highly complex to use and can be not very user-friendly unless you are already knowledgeable of the subject. Todo

> *Explain what is topic malfunctioning.* <

ROSMonitoring also doesn't provide a ROS-oriented DSL, instead, it takes advantage of the already built and somewhat generic language Runtime Monitoring Language (RML) [1] which doesn't provide the expressiveness and intuitiveness I expect from a ROS-oriented invariant specification language.

ROSRV [8] is another tool that provides the runtime monitoring of ROS systems, however, ROSRV implements safety measures while the system is running, interfering with its normal execution. ROSRV lacks documentation on its functioning and usage. Additionally, ROSRV seems to have been discontinued since its last version is only capable of running under ROS Groovy Galapagos, which has not been supported since 2014.

Chapter 3

A Specification Language for Robotics Properties

In this chapter, the structure and intricacies of the DSL are presented. The notations used in the DSL, like concepts and keywords, are introduced in (section 3.1). The DSL grammar is written in the Backus-Naur Form (BNF) (section 3.2). BNF is a notation used to describe the grammar or syntax of computable languages. Finally, some practical examples are written with the help of the DSL to display its expressiveness (section 3.3).

3.1 Language Notations

Todo >Introduce the following subsections.<

3.1.1 High-level Concepts

The high-level concepts that can be created in the language are:

- **Property** - A property represents a temporal specification or a blend of temporal specifications between components.
- **Declaration** - A declaration allows for the representation of ROS *topics* in order to interact with them.
- **Model** - A model allows for the declaration of specific *topics* that are required when correlating certain robots' and simulation components.
- **Association** - An association serves as a way to create program variables.

3.1.2 Temporal Keywords

The language considers not only LTL basic operators but also some common shortcuts for useful combinations of such operators, like *after_until*.

- **always X** - X has to hold on the entire subsequent path;
- **never X** - X never holds on the entire subsequent path;

- **eventually X** - X eventually has to hold somewhere on the subsequent path;
- **after X, Y** - after the event X is observed, Y has to hold on the entire subsequent path;
- **until X, Y** - X holds at the current or future position, and Y has to hold until that position. At that position, Y does not have to hold anymore;
- **after_until X, Y, Z** - after the event X is observed, Z has to hold on the entire subsequent path up until Y happens. At that position, Z does not have to hold anymore;

3.1.3 Temporal value

Todo *>X is a property before. Now it is a variable. I would recommend using lowercase x for variable names, or other letter.<*

It is also possible to reference previous variable states:

$$@\{X, -y\} \tag{3.1}$$

This will represent the value of the variable X in the point in time -y.

3.1.4 Simulation primitives

To support comparing the internal state of the robotic system with the environment, the language provides basic primitives to refer to the simulation environment:

- **X.position** - The position of the robot in the simulation;
- **X.position.y** - The position in the y axis of the robot in the simulation. Also works for x and z;
- **X.distance.Y** - The absolute distance between two objects in the simulation. For the x and y axis;
- **X.distanceZ.Y** - The absolute distance between two objects in the simulation. For the x, y, and z axis;
- **X.velocity** - The velocity of an object in the simulation. This refers to linear velocity;
- **X.velocity.x** - The velocity in the x axis of an object in the simulation. This refers to linear velocity;
- **X.localization_error** - The difference between the robot's perception of its position and the actual position in the simulation;

3.1.5 Operands

Besides the already mentioned operands, *Temporal values*, *Simulation primitives*, and *Temporal Keywords*, the DSL also supports both Integer and Float values, Booleans, and declared variables.

3.1.6 Operators

The DSL supports operators to correlate components. The operators are $+$ (addition), $-$ (subtraction), $*$ (multiplication), $/$ (division), $==$ (equals), $!=$ (different), $>$ (greater than), $>=$ (greater or equal than), $<$ (lower than), $<=$ (lower or equal than), *and* (conjunction), *or* (disjunction), *implies* (implication), and for any comparison operator $X \ Xy$ - the values being compared will have an error margin of y (Example: $Z == 0.05 \ Y$).

3.1.7 Protected Variables

Protected variables are variable names restricted to set determined monitoring parameters.

`__rate__` - Set the frame rate which properties are checked (By default, the rate is 30hz)

`__timeout__` - Set the timeout for how long the verification will last (By default, the timeout is 100 seconds)

`__margin__` - Set the error margin for comparisons

3.1.8 Topic declaration

In order to relate robot components with the simulation, the developer can declare the relevant *topics*.

The language cannot inherently have a way to interact with specific components of a robot because it does not know which *topic* to get information from. Therefore, the developer needs to declare these specific topics to be able to interact with them.

Todo *>Usar listings nos exemplos seguintes, como aqui.<*

```
1 decl robot_position /odom Odometry.pose.pose.position
```

The variable robot_position was declared with the type Odometry.pose.pose.position and is linked to the topic /odom

3.1.9 Model robots

A set of specific topics can be modeled for the robot, like *position* or *velocity*. The compiler will use these to call specific functions that need this information from the robot's perspective.

```
1 model robot1:
2   position /odom Odometry.pose.pose.position
3   ;
4   never robot1.localization\_error > 0.002
```

The localization_error function requires the position topic of a robot to be modeled in order to compare it with the actual simulation position

3.2 Grammar

Todo *>you need a sentence introducing the grammar. You should also highlight the need for a few productions, maybe relating with the high-level concepts.<*

| | | |
|------------------|-----|---|
| <program> | ::= | <command> <command> <program> |
| <command> | ::= | <association> <declaration> <model> <pattern> |
| <association> | ::= | name = <pattern> _rate_ = integer _timeout_ = <number> _default_margin_ = <number> |
| <declaration> | ::= | decl name topic_name <msgtype> decl name name <msgtype> |
| <model> | ::= | model name : <modelargs> ; |
| <modelargs> | ::= | <name> topic_name <msgtype> <name> <name> <msgtype> <name> topic_name <msgtype> <modelargs> <name> <name> <msgtype> <modelargs> |
| <name> | ::= | name <func_main> |
| <func_main> | ::= | position velocity distance localization_error orientation |
| <msgtype> | ::= | <name> <name> . <msgtype> |
| <pattern> | ::= | always <pattern> never <pattern> eventually <pattern> after <pattern> , <pattern> until <pattern> , <pattern> after_until <pattern> , <pattern> , <pattern> <conjunction> |
| <conjunction> | ::= | <conjunction> and <comparison> <conjunction> or <comparison> <conjunction> implies <comparison> <comparison> |
| <comparison> | ::= | <multiplication> <opbin> <multiplication> <multiplication> <opbin> <number> <multiplication> <multiplication> |
| <opbin> | ::= | < > <= >= == != |
| <multiplication> | ::= | <multiplication> * <addition> <multiplication> / <addition> <addition> |
| <addition> | ::= | <addition> + <operand> <addition> - <operand> <operand> |
| <operand> | ::= | name <number> true false <func> <temporalvalue> (<pattern>) |
| <number> | ::= | float integer |
| <func> | ::= | name . <func_main> name . <func_main> <funcargs> |
| <funcargs> | ::= | . <name> . <name> <funcargs> |
| <temporalvalue> | ::= | @ name , integer |

3.3 DSL Usage Examples

To validate the expressive power of our language, I present some examples of expressions inspired by real-world scenarios.

3.3.1 Vehicle Maximum Speed

Some robots have a maximum safe speed at which they can move. Sometimes this limit is imposed by law, but some other times by physical constraints.

The robot velocity will never be above 2 for the duration of the simulation;
never robot.velocity > 2.0

3.3.2 Follow the Leader

The first robot being above 1 velocity implies that the second robot is at least at 0.8 distance from the first robot, up until the first robot reaches a particular location;

until (robot1.position.x > 45 and robot1.position.y > 45), always (robot1.velocity > 1 implies robot2.distance.robot1 > 0.8)

3.3.3 Localization error

The localization error (difference between the robot's perception of its location and the actual simulation location) of the robot is never above a specific value.

model robot1:
 position /odom Odometry.pose.pose.position
 ;
 never robot1.localization_error > 0.002

3.3.4 Drone height rotors control

After a drone is at a certain altitude, both rotors always have the same velocity up until the drone decreases to a certain altitude.

decl rotor1_vel /drone_mov/rotor1 Vector3.linear.x
decl rotor2_vel /drone_mov/rotor2 Vector3.linear.x
after_until drone.position.z > 5, drone.position.z < 5, rotor1_vel == rotor2_vel

Chapter 4

Monitoring

This chapter explains the whole process of monitoring, from compilation to error detection. First, the overall process of compilation is explained (section 4.1), then the generated file and some of its specifications are described (section 4.2), and finally, representative error messages are presented (section 4.3).

4.1 Runtime Monitoring

After writing all the desired robotic systems specifications, the file needs to be compiled to a monitoring Python module. This is currently done running the following script, from its location: `python language.py properties.txt /home/ros_workspace/src/my_pkg/src`.

The *language.py* file needs to be run as a Python file and takes as arguments:

1. The specifications file;
2. The path where the the generated Python monitoring module will be generated.

The given directory for the generated file should be under a ROS workspace for the compilation to succeed. This is because, during the compilation, access to information like the available ROS messages might be necessary.

The monitoring file can now run as an independent ROS node, integrated into a launch file, or using *roslaunch* in the console to execute it.

4.2 Generated File

Todo >New name of the section: *Compilation*<

Todo >This section needs to be vastly extended! It should have 1) figure with architecture of the generated model; 2) A description of the runtime library you have to support the property checking (i.e., auxiliary functions) and 3) The translation rules between the source code and the generated code. Examples of things that need: how to subscribe to topics that were declared, how to emit error messages, how to access past versions of the state, how is it recorded (data structure).<

The generated monitoring file declares the needed subscribers and uses `ApproximateTimeSynchronizer` to call the callback function. The `ApproximateTimeSynchronizer` synchronizes messages by their timestamp and if they do not have a header, use the ROS time.

A callback function is called every time a new message from one of the subscribers is received. The callback function saves the relevant information for property checking in a global variable. This information serves as a current "screenshot" of the simulation representing its current state.

The node executes a loop at a delineated rate, doing the following tasks:

1. Check if the defined simulation timeout time has reached.
2. Save the current simulation state obtained by the callback function. This is necessary because the callback function is called at fluctuating rates, and the objective is to save multiple "screenshots" of the simulation at the loop fixed rate to make correlations with past states.
3. Verify the properties using the saved states and calling each created function. An independent function with the necessary computations for verifying the property is defined for each base property.

4.3 Error Messages

An error message starts by stating the line in the specification file which resulted in an error and showing the specification itself. **Todo** *>This does not make sense as the explanation for error messages! You have different temporal restrictions: are the errors the same for all of them? How do you select the state that you present? Do you show all the variables in context, the monitored ones, or the ones present in the property? There is space for discussion here on how you could provide more precise messages.<*

Afterward, the value at the time of failure of all the variables present at the specification is shown.

Todo *>Lots of whitespace.<*


```
Error at line 19:  
until turtlebot3_burger.position.x <= -1 and turtlebot3_burger.position.y < -1,  
never turtlebot3_burger.velocity >= 0.1  
Failing state:  
  turtlebot3_burger.position.x: 0.1134222404473394  
  turtlebot3_burger.position.y: 0.0011127060961216948  
  turtlebot3_burger.velocity: 0.10048210526931797
```

Figure 4.1: Example of an error message.

Chapter 5

Evaluation

This chapter introduces the evaluation process of the work. [section 5.1](#) gives a broad overview of the whole process, and [section 5.2](#) goes into more detail about each one of the experiments.

5.1 Evaluation Overview

To evaluate the developed work, I decided to go through a list of already documented ROS bugs and identify three that happen at runtime. After that, I specified a robot's properties in the DSL that should be capable of detecting an error for said bugs while running the system.

ROBUST [2] is a dataset that documents over two hundred bugs in multiple robots using ROS. After going through the dataset, three bugs were chosen:

- Calculation Error Inverts Turning Direction ¹ ([subsection 5.2.1](#)) - "Due to an error in velocity calculations, when Kobuki was issued a very low negative linear speed (very slow backwards movement), it would also inadvertently invert its turning direction. That is, if it was supposed to move backwards while turning left, it would move backwards and turn right instead."
- Robot Getting Stuck When Auto-docking ² ([subsection 5.2.2](#)) - "The movement speeds were hard-coded for the auto-docking algorithm, and worked well for regular Kobuki and Turtlebot, but were too slow for heavier robots, causing them to get stuck."
- Unexpected Movement Due to Wrong Calculation ³ ([subsection 5.2.3](#)) - "Kobuki moves using differential drive. Originally, the command velocities (linear and angular) were provided as 'short', and were converted to 'short' after each step, even though the calculations yielded floating point numbers. This lead to calculation errors in some special cases, where the robot was supposed to move forward but ended up moving backwards instead."

To replicate each bug, a *daemon* node is inserted into the system that interferes with the normal runtime flow and replicates the desired bug. [Figure 5.1](#) represents the natural flow of the systems whilst [Figure 5.2](#) shows the system flow when trying to replicate a bug with the help of the *daemon* node.

¹<https://github.com/robust-robin/robust/blob/master/kobuki/e964bbb/e964bbb.bug>

²<https://github.com/robust-robin/robust/blob/master/kobuki/0416c81/0416c81.bug>

³<https://github.com/robust-robin/robust/blob/master/kobuki/1c141a5/1c141a5.bug>




Figure 5.1: The normal runtime flow of the system.




Figure 5.2: The flow of the system when adding the daemon node.

When injecting a bug, all the data before addressed to the robots' `/cmd_vel` topic is remapped to a new topic called `/cmd_vel_original`, the `daemon` node subscribes to this new topic and modifies the data before sending it to the robots' `/cmd_vel` topic so that the robot behaves like the expected bug.

With this new system flow, it is possible to specify properties between the `/cmd_vel` topic which represents the robots' behavior, and the `/cmd_vel_original` topic, which represents the actual command given to the robot.

5.2 Experiments

This section goes through the property specification and runtime monitoring for the three mentioned selected bugs. Calculation Error Inverts Turning Direction (section 2.1), Robot Getting Stuck When Auto-docking(section 2.1), and Unexpected Movement Due to Wrong Calculation (section 2.1).

5.2.1 Calculation Error Inverts Turning Direction

DSL property specification:

First, I declare the `cmd_vel_original` topic, which will represent the commands given to the robot. Then I make a correlation between the topic that represents the given commands and the topic that represents the actual robot's behavior.

```

decl angular_vel_robot_perception cmd_vel_original Twist.angular.z
after turtlebot3.velocity.angular.z < 0, never angular_vel_robot_perception > 0

```

Translating into natural language, the property states in the first section that after the robot's actual simulation z parameter of the angular velocity is less than zero, then in the second section, never the z parameter of the angular velocity of the command given to the robot is more than zero.

Daemon node code and behavior:

```
Error at line 74:
  after turtlebot3.velocity.angular.z < 0, never a
  ngular_vel_robot_perception > 0
Failing state:
  turtlebot3.velocity.angular.z: -0.00144950743
527
```

Figure 5.3: Calculation Error Inverts Turning Direction bug error message.

```
1 class Direction_invert_error:
2
3     def __init__(self):
4         print("simulating direction_invert_error_behavior...")
5         self.cmd_vel_pub = rospy.Publisher("cmd_vel", Twist, queue_size=1)
6         self.twist = Twist()
7         self.direction_invert_error()
8
9     def get_vel(self):
10        return rospy.wait_for_message("cmd_vel_original", Twist)
11
12    def direction_invert_error(self):
13        while not rospy.is_shutdown():
14            vel = self.get_vel()
15            self.twist = vel
16            if abs(vel.linear.x) < 0.012:
17                self.twist.angular.z = -vel.angular.z
18            self.cmd_vel_pub.publish(self.twist)
```

The *daemon* node checks when the given robot's command linear velocity is below *0.012* and injects the opposite value of the *z* parameter of the angular velocity to the actual robot's velocity in the simulation.

Now when giving commands to the robot, if the given velocity is below *0.012* and I make a turn, the robot will turn the opposite way. The output of the monitoring node for a test case is demonstrated in [Figure 5.3](#).

5.2.2 Robot Getting Stuck When Auto-docking

DSL property specification:

First, I declare the `cmd_vel_original` topic, which will represent the commands given to the robot. Then I make a correlation between the topic that represents the given commands and the topic that represents the actual robot's behavior.

```
decl vel_robot_perception cmd_vel_original Twist.linear.x
after vel_robot_perception > 0, never turtlebot3.velocity == 0.005 0
```

*Translating into natural language, the property states in the first section that after the given robot's commands *x* parameter of the linear velocity is greater than zero, then in the second section never, the actual simulation's linear velocity is equal to zero.*

Daemon node code and behavior:

```
1 class Auto_docking_error:
2
3     def __init__(self):
4         print("simulating auto_docking_error_behavior...")
```

```

Error at line 78:
after vel_robot_perception > 0, never turtlebot3.velocity =={0
.005} 0
Failing state:
turtlebot3.velocity: 3.59827874925e-05

```

Figure 5.4: Robot Getting Stuck When Auto-docking bug error message.

```

5     self.cmd_vel_pub = rospy.Publisher("cmd_vel", Twist, queue_size=1)
6     self.twist = Twist()
7     self.auto_docking_error()
8
9     def get_vel(self):
10         return rospy.wait_for_message("cmd_vel_original", Twist)
11
12     def auto_docking_error(self):
13         while not rospy.is_shutdown():
14             vel = self.get_vel()
15             self.twist = vel
16             if abs(vel.linear.x) < 0.015:
17                 self.twist.linear.x = 0.0
18             self.cmd_vel_pub.publish(self.twist)

```

The *daemon* node checks when the given robot's command linear velocity is below 0.015 and injects a value of 0.0 to the linear velocity of the actual robot's velocity in the simulation.

Now when giving commands to the robot, if the given velocity is below 0.015 the robot will stay stationary. The output of the monitoring node for a test case is demonstrated in [Figure 5.4](#).

5.2.3 Unexpected Movement Due to Wrong Calculation

DSL property specification:

First, I declare the `cmd_vel_original` topic, which will represent the commands given to the robot. Then I make a correlation between the topic that represents the given commands and the topic that represents the actual robot's behavior.

```

decl vel_robot_perception cmd_vel_original Twist.linear.x
after turtlebot3.velocity.linear.x < 0, never vel_robot_perception > 0

```

Translating into natural language, the property states in the first section that after the robot's actual simulation x parameter of the linear velocity is less than zero, then in the second section, never the x parameter of the linear velocity of the command given to the robot is more than zero.

Daemon node code and behavior:

```

1     class Backwards_movement_error:
2
3         def __init__(self):
4             print("simulating backwards_movement_error_behavior...")
5             self.cmd_vel_pub = rospy.Publisher("cmd_vel", Twist, queue_size=1)
6             self.twist = Twist()
7             self.backwards_movement_error()
8
9         def get_vel(self):
10             return rospy.wait_for_message("cmd_vel_original", Twist)
11
12         def backwards_movement_error(self):
13             while not rospy.is_shutdown():

```

```
Error at line 82:  
after turtlebot3.velocity.linear.x < 0, never vel_robot_percep  
tion > 0  
Failing state:  
turtlebot3.velocity.linear.x: -0.0387927308907
```

Figure 5.5: Unexpected Movement Due to Wrong Calculation bug error message.

```
14         vel = self.get_vel()  
15         self.twist = vel  
16         if abs(vel.linear.x) > 0.03:  
17             self.twist.linear.x = -vel.linear.x  
18             self.cmd_vel_pub.publish(self.twist)
```

The *daemon* node checks when the given robot's command linear velocity is above 0.03 and injects the opposite value of the x parameter of the linear velocity to the actual robot's velocity in the simulation.

Now when giving commands to the robot, if the given velocity is above 0.03 , the robot will start moving backward. The output of the monitoring node for a test case is demonstrated in [Figure 5.5](#).

Todo >Feels like it's missing a section with a list of real-world bugs that could be detected with your language, even if you did not implement them.<

Chapter 6

Future Work

In this chapter the possible work left undone or that could improve our study is presented. Improvement of the developed work performance is mentioned in [section 6.1](#), [section 6.2](#) mentions the validation of the work, [section 6.3](#) discusses about the works' error messages, [section 6.4](#) talks about the integration with scenario generation tools, and finally [section 6.5](#) mentions the possible integration with other simulators besides Gazebo.

6.1 Performance Tweaking

The generated code performance can be improved so that the load of the monitoring node on the simulation is reduced.

For instance, the frequency at which some properties are checked could fluctuate. In some circumstances, a particular property does not need to be checked at every simulation iteration. Implementing some mechanism that can skip certain property checks per iteration will undoubtedly decrease the load the monitoring node will have on the simulation.

6.2 Validation of the Proposal

Although some evaluation was done for the work done, more evidence and experimental data on the effective capabilities of the proposal are still needed:

1. How expressive is the DSL from the developers' point of view in specifying robots' properties.
2. Proof of concept that the system is able to detect the rule violations specified by the DSL.
3. Evidence that the monitoring does not disturb the simulation by demanding excessive resources.

6.3 Better Error Messages

Giving developers helpful and comprehensible error messages is a shared concern amongst all compilers. One can argue that even the best compilers still have space for improvement when talking about delivering good error messages.

Although in our work I gave some thought to the error messages delivery, I believe that a more narrow error localization is still possible.

Also, there was no time for a thorough validation of the proposal, which means that some bugs could be present when delivering the error messages, and the users could have some problems with the delivery or ideas on how to improve it.

6.4 Integrate the DSL with Scenario Generation Tools

Integrating our work with a scenario generation tool would improve the whole test automation process by creating unpredictable environments on where to test our specified properties, allowing the execution of multiple tests.

For instance, GzScenic [4] is a tool that generates random scenarios for the Gazebo simulator based on a defined model.

6.5 Integration with other Industrial Simulators

Our work was developed with the Gazebo simulator in mind, which means the monitoring is currently not compatible with other simulators.

Although Gazebo is widely used, many other simulators are also currently being used. Therefore, adapting our work to integrate other widely used simulators would be helpful for many users that choose not to use Gazebo.

Todo >What other simulators are there?<

Todo >This work can be used to measure quality of software modules. thus, it can be used in automated program synthesis, repair and improvement<

Chapter 7

Conclusion

Due to the fact that robots interact with the real world, robotic systems are unpredictable. Coming up with a reliable and efficient method for automatic robot testing is a challenge, one of the reasons being that verifying the success of a task may not be possible from the robot's perspective. Current practices in testing robot software mainly require a human to analyze the robot's behavior to determine its correctness, one way of overcoming this problem is with some other type of automatic external monitoring.

Our approach relies on simulation-based testing so that developers can take advantage of the real values of objects' attributes on the simulation to compare with what the robot system perceives, trying in this way to surpass the need for human-in-the-loop testing.

I developed a DSL that allows for the specification of a robotic system's properties, designed from the ROS framework point of view, and that abstracts the underlying LTL system. As a result, it is possible to express relevant temporal and positional arguments between robots' components and objects in the simulation and also properties that relate the internal information of the system with the corresponding information in the Gazebo simulator.

I automated the generation of a monitoring Python file that can monitor some behavioral violations of robots in relation to their state or events during a Gazebo simulation.

I have shown that the approach can monitor some interesting scenarios that developers care about. I also present what is still left to do and what future work is still needed.

Todo > *This seems more like a summary and not a conclusion. A conclusion could focus more on the results obtained, and how that changes the world.* <

Bibliography

- [1] Runtime Monitoring Language. <https://rmlatdibris.github.io/>. [Online; accessed 13-September-2022].
- [2] ROBUST: ROS Bug Study. <https://github.com/robust-rocin/robust>. [Online; accessed 28-August-2022].
- [3] ROS-INDUSTRIAL. <https://rosindustrial.org/>. [Online; accessed 11-September-2022].
- [4] Afsoon Afzal, Claire Le Goues, and Christopher S. Timperley. GzScenic: Automatic scene generation for gazebo simulator, 2021.
- [5] Afsoon Afzal, Claire Le Goues, and Christopher Steven Timperley. Mithra: Anomaly detection as an oracle for cyberphysical systems. *IEEE Transactions on Software Engineering*, pages 1–1, 2021. doi: 10.1109/TSE.2021.3120680.
- [6] Matthew B Dwyer, George S Avrunin, and James C Corbett. Property specification patterns for finite-state verification. In *Proceedings of the second workshop on Formal methods in software practice*, pages 7–15, 1998.
- [7] Angelo Ferrando, Rafael C Cardoso, Michael Fisher, Davide Ancona, Luca Franceschini, and Viviana Mascardi. Rosmonitoring: a runtime verification framework for ros. In *Annual Conference Towards Autonomous Robotic Systems*, pages 387–399. Springer, 2020.
- [8] Jeff Huang, Cansu Erdogan, Yi Zhang, Brandon Moore, Qingzhou Luo, Aravind Sundaresan, and Grigore Rosu. Rosrv: Runtime verification for robots. In *International Conference on Runtime Verification*, pages 247–254. Springer, 2014.
- [9] Nathan Koenig and Andrew Howard. Design and use paradigms for gazebo, an open-source multi-robot simulator. In *2004 IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)(IEEE Cat. No. 04CH37566)*, volume 3, pages 2149–2154. IEEE, 2004.
- [10] Sophia Kolak, Afsoon Afzal, Claire Le Goues, Michael Hilton, and Christopher Steven Timperley. It takes a village to build a robot: An empirical study of the ros ecosystem. In *2020 IEEE International Conference on Software Maintenance and Evolution (ICSME)*, pages 430–440, 2020. doi: 10.1109/ICSME46990.2020.00048.
- [11] Amir Pnueli. The temporal logic of programs. In *18th Annual Symposium on Foundations of Computer Science (sfcs 1977)*, pages 46–57. iee, 1977.

- [12] Morgan Quigley, Ken Conle, Brian Gerkey, Josh Faust, Tully Foote, Jeremy Leibs, Rob Wheeler, and Andrew Ng. Ros: an open-source robot operating system. *ICRA Workshop on Open Source Software*, 3(3.2):1–6, 01 2009.
- [13] Marco Stadler, Michael Vierhauser, and Jane Cleland-Huang. Towards flexible runtime monitoring support for ros-based applications. In *RoSE'22: 4th International Workshop on Robotics Software Engineering Proceedings*, 2022.
- [14] Milda Zizyte, Casidhe Hutchison, Raewyn Duvall, Claire Le Goues, and Philip Koopman. The importance of safety invariants in robustness testing autonomy systems. In *2021 51st Annual IEEE/IFIP International Conference on Dependable Systems and Networks-Supplemental Volume (DSN-S)*, pages 41–44. IEEE, 2021.